A Machine Learning Framework for Detecting Refund Abuse in E-Commerce Websites

Devansh Mane
Computer Science and Engineering (AI and ML)
A.P. Shah Institute of Technology
Thane, India
devanshmane1604@gmail.com

Vedant Gandre
Computer Science and Engineering (AI and ML)
A.P. Shah Institute of Technology
Thane, India
vedantgandre11@apsit.edu.in

Harshit Jain
Computer Science and Engineering (AI and ML)
A.P. Shah Institute of Technology
Thane, India
harshitjain39@apsit.edu.in

Sumeet Shingi
Computer Science and Engineering (AI and ML)
A.P. Shah Institute of Technology
Thane, India
svshingi@apsit.edu.in

***Abstract*— Refund abuse is becoming a real issue in online shopping, and it is also causing financial losses across the e-commerce ecosystem. These systematic returns often slip through the basic security setups that companies have. This paper brings up Fraud Lens, which is basically a machine learning framework to catch fraudulent return habits at an early stage. It tries to do that while still maintaining a positive user experience for regular users. Derived from recent studies on how people behave in e-commerce and those integrated assistant tools is appropriate for this context. The framework works through a multi-stage pipeline that digs into huge amounts of transaction data. The key part here is a dual-purpose engine. First, a predictive model is employed that figures out a specific Return Rate Percentage for every user, enabling the identification of the high-risk ones from the legitimate customers. Second, it adds in a module for comparing prices across platforms. That turns the whole fraud tool into something useful like a shopping assistant. Users are provided with aggregated with real-time market deals, which adds value. At the same time, it gives a baseline for spotting price-sensitive patterns in shopping behaviour. It also contributes to fraud detection. They evaluated different ML models comparatively and used Natural Language Processing to generate narratives about fraud automatically. This reduces manual investigation efforts and reduces operational latency. From the experiments, results indicate that, connecting revenue protection with features that put consumers first. Overall, this seems like a scalable fix that fits into the trend toward more vigilant e-shopping environments.**

Keywords—Machine Learning, Refund Abuse, E-Commerce, Return Rate Percentage, Price Comparison, Fraud Detection, Behavioural Analytics.

I. Introduction

E-commerce has experienced rapid growth in recent years, where everyone adopting online purchasing. Billions of people are dependent on digital transactions and so it has become the important mode of commerce But along with that, fraud has also seen a huge crime rate in recent years, and is growing in frequency and sophistication. Most research before looked at third-party stuff like stealing credit cards or taking over accounts. Still, there is one another challenge that majorly impacts business profitability, is the refund abuse, which is commonly referred to as first-party fraud.

It happens when real customers take advantage of easy return rules. Like purchasing clothes for temporary use followed by return, it is also known as wardrobing, or saying the package never delivered so they get money back but keep the item. According to [2], looking at returns is not just about logistics anymore, its key for security too, because it can hide bigger patterns of taking advantage.

Usual fraud systems miss this since everything looks normal when they buy. As [10] points out, those old rule-based setups treat everything the same, but fraudsters keep changing how they do things. So, we really need smarter automated ways to distinguish loyal customers from people who return too much.

This paper brings in Fraud Lens; it's a machine learning based framework to comprehensive view if customers are reliable. Other models just check the one purchase [3] [6], but this looks at the whole time someone uses the site.

It has two main parts. First, risk quantification, where it figures out a return rate percentage that always update, so admins get a score to decide if an account needs watching or blocking. Second, its user-friendly, like the E-Vigilant model in [11], with a price comparison tool that shows real-time deals to help people save, and it also tracks if low prices link to more returns, which may reveal behavioural patterns.

The idea here is to go from just spotting fraud after it happens to something more active with analysis. Contributions include building an ML setup that finds weird return habits in big data sets. Then there is the price feature for keeping users around and watching behaviour. Also, using NLP to generate narratives that are interpretable from model outputs, so investigators can understand easier.

Fraud Lens shows platforms can have tight security without making customers mad, through smart tech achieving this balance remains challenging but it might work despite potential concerns regarding simplification.

II. Literature Review

Despite effortless access to online stores, weaknesses in their infrastructure grow just as fast. Research shows safeguards have outgrown basic login methods long ago. Now attention turns to behavioral patterns hidden within customer information. Instead of passwords, complex algorithms detect irregularities by analyzing chaotic flows of purchase records. These automated tools handle volumes of activity no

XXX-X-XXXX-XXXX-X/XX/$XX.00 ©20XX IEEE

team could review by hand. Harder choices emerge when convenience meets unseen risks beneath the surface.

In practice, movement within the field now follows patterns labeled “multi-stage pipelines” by specialists [3]. Instead of viewing each transaction separately, such systems search for recurring behavioral traces. Through exposure to vast amounts of data, these algorithms begin distinguishing simple errors from deliberate manipulation. Such precision matters deeply; research shows overly strict controls may falsely label activity as fraudulent - a misstep that can alienate legitimate users [5][13]. Success means more than correctness - it includes clarity, offering explanations clear enough for human review [6].

A different pattern now shapes how such systems engage users. Not only do certain applications defend organizational interests, yet they simultaneously support individual needs. Such designs sometimes carry the label "E-Vigilant" [11]. Where some analysts apply image-based methods to reduce unintentional product returns [4], alternative studies probe patterns behind deliberate return actions [2]. This is where monetary losses accumulate silently - the pattern often termed the "hidden bleed" within online retail. A noticeable absence exists in present research regarding systems capable of identifying habitual returners prior to profit erosion [12].

Yet transparency fades when machines gain influence, raising concerns about hidden decision paths. Because clarity matters, scientists now probe deeper into systems that reveal their reasoning. A fraud tracker might clarify risks by generating simple reports - telling staff what it sees without jargon [7]. Such openness supports fairness, reducing chances the software favors one group over another unknowingly [8][14].

Lately, efforts have taken shape to bring consistency to risk measurement methods. Increasingly, structures such as Graph Neural Networks uncover concealed fraud groups [7], while indicators including Return Rate Percentage deliver straightforward numerical evaluations of exposure [9]. According to recent findings, agreement has formed: artificial intelligence excels at detecting trends - yet proves most valuable when guiding decisions made by individuals, not substituting them. Systems showing greatest impact, comparable to the model described here, transform intricate information into understandable conclusions aimed at key users [10].

III. METHODOLOGY

Despite common approaches, this tool targets first-party fraud through layered modules within a web environment. Originating from Flask's lightweight structure, the setup combines live data collection with detailed behavior tracking across user sessions. Notably, decisions emerge not from isolated events but patterns drawn over time, guided by methods referenced in [3] and refined using comparison strategies from [5]. Instead of relying solely on pass-or-fail judgments, the approach measures trustworthiness along shifting scales. Central to its function, one component computes precise return likelihoods while simultaneously supporting an interactive guide at the interface level - a feature shaped partially by concepts in [11]. Rare among similar tools, it balances enforcement mechanisms with usability considerations without favoring one over the other. This method uses two layers so the system tracks deliberate return inconsistencies [2], yet delivers practical tools such as pricing checks across platforms. Since XGBoost works alongside Neural Networks, detection of irregular fraud behaviors becomes possible; at the same time, a rule-driven NLP component converts intricate results into clear written summaries. As a result, transparency is maintained, meaning explanations stay accessible to administrative staff.

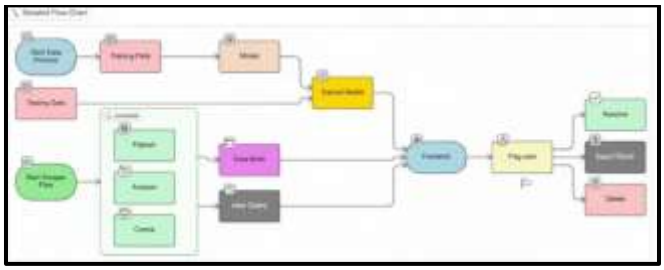

Figure1: Hybrid Machine Learning Model Combining XGBoost and Neural Network

Figure 1 displays how Fraud Lens operates as a whole. Starting simultaneously, one path handles past records to refine combined prediction methods. Meanwhile, another gathers current marketplace details from sites such as Amazon and Flipkart without delay. These paths meet inside the Data Node prior to appearing on the interface. When a request arrives, evaluation begins immediately. Should irregularities appear, alerts activate under the Flag User procedure. Administrators then receive choices - export results into spreadsheet format, address concerns directly, or remove questionable entries entirely.

A. Core System Architecture and Technologies

At its core lies a unified structure formed by powerful tools from the open community. This framework combines components designed for speed through collaborative development efforts. Built without proprietary restrictions, it relies on transparent systems working together seamlessly. Performance emerges from carefully linked elements sharing common standards. Each part connects to support reliability across diverse environments. Foundation strength comes from consistent interaction among independent yet interdependent modules

The core structure relies on Flask, directing RESTful API calls while linking machine learning components. From session tracking to past transactions, data organization runs through Flask-SQLAlchemy. A PostgreSQL system holds the information, arranged with precision. Integration unfolds behind each interaction, ensuring consistency across operations.

A Machine Learning Engine operates through two components working together. One part applies XGBoost to manage boosted decision trees. Meanwhile, the other uses a Sequential Neural Network - developed with Keras and TensorFlow - for capturing complex feature patterns. This structure allows distinct methods to address different aspects of prediction tasks.

From time to time, live pricing details are collected through automated methods. One tool processes website layout by interpreting raw markup elements directly. At the same time, another handles communication across multiple networks requests efficiently. Together, they enable continuous monitoring without manual involvement. These components operate independently but align closely during execution cycles. Their integration supports consistent access to updated

figures online. Each part performs under variable conditions while maintaining accuracy overall. Efforts rely on structured navigation rather than visual cues throughout. Despite differing roles, synchronization remains stable even at scale. The method adapts quickly when sources modify their output format unexpectedly. Throughout operation, minimal overhead is preserved systematically. Such design allows reliable extraction despite frequent site updates elsewhere.

With security in mind, customer identifiers appear only as hashed values. Though relations matter, the structure links Customer_ID to Order_History and Referral_Flags without gaps. From stability arises consistency - data ties remain intact by design

B. Data Acquisition and Knowledge Base Creation

One kind of information enters through a dedicated channel, while another follows a separate path entirely. Though both are handled by the same framework, their entry methods differ in structure and timing. A first stream flows under specific conditions, whereas the second activates only when predefined triggers occur. Each route maintains its own rules, yet they converge later within the core mechanism

1. Internal Fraud Information Collection Transactions

Originating from archived records of online purchases, the dataset includes details on buyer information along with patterns in product returns. Stored systematically, it reflects activity across multiple seasons without real-time updates. Entries accumulate through automated logging mechanisms rather than manual input. Information flows into the structure after verification steps ensure consistency over time. Each record forms part of a larger sequence shaped by past interactions. Built on repeated cycles of purchase and feedback, the collection grows incrementally behind secure access layers.

A sequence of steps transforms raw inputs before analysis begins. Through this pipeline, numerical values undergo scaling adjustments. Categorical fields experience conversion via encoded representation methods. Important characteristics emerge during processing - examples include how often returns occur and typical item cost levels. These outputs become inputs for machine learning structures.

Because fraud occurs less frequently than valid activity, synthetic data generation adjusts proportions across categories. This method prevents skewed assumptions during analysis by maintaining equal representation. Balance emerges through structured replication rather than natural occurrence alone.

2. Market Intelligence Base Price Corpus

Data Source Real time pricing data from Amazon eBay Flipkart and Myntra.

A dedicated scraping process links each product SKU to multiple online platforms. Through BS4, extraction occurs for Price Tag, alongside Currency and Availability Status. Stored temporarily in Redis, the collected information waits. Instant access is enabled for the price comparison component. Repeated scrapes of identical items pause when under thirty minutes have passed since last fetch.

A. Module-Specific Implementation

1. Hybrid Refund Abuse Detection

Beginning with structured data - such as rates of returned orders - an XGBoost model initiates the process. Following that phase, patterns across time-stamped interactions are examined by a neural architecture. Rather than relying on one method alone, outcomes merge through weighted voting principles. When likelihood scores surpass preset thresholds, human evaluators receive notification automatically. Such design emerges from combining distinct analytical strengths sequentially.

2. Risk Quantification Module the RRP Engine

The Return Rate Percentage Is Calculated Per Customer ID Using Mathematical Logic

Whenever (RRP) exceeds 40%, a critical alert activates. Because of this threshold, the platform detects frequent returners without flagging rare cases among typical users.

$$RRP = \left(\frac{\sum Returned_Items}{\sum Total_Orders}\right) \times 100$$

3. Cross Platform Price Comparison Tool

This function supports both practical access and tracking of actions. Through FuzzyWuzzy comparison, matches form between items listed on different sites. A sorted overview of favorable options appears for the person, whereas internal records capture how pricing influences choices, shaping a clearer picture of individual decision patterns over time. Despite appearances, subtle shifts in response guide long-term adjustments behind the scenes.

IV. RESULTS

An Experimental Setup Environment

One step at a time, the Fraud Lens system unfolded inside a test setting built around today's split-architecture AI models. A desktop running Python 3.10 guided its construction - powered by an Intel i7 chip, supported with 16 gigabytes of memory. From underneath, Flask held things together, stitching each part into place without force or fuss

System Components and Key Tools

- Flask Backend with PostgreSQL and User Authentication
- Machine Learning Models: XGBoost and TensorFlow Keras Sequential Neural Network
- Data Processing with Scikit-learn SMOTE Min Max Scaler Pandas NumPy
- Scraping Engine BS4 Requests Redis Caching
- NLP Reporting with Jinja2 Templates and NLTK Text Processing
- HTML5 CSS3 JavaScript with AJAX for live data fetching
- B System Interface and Module Outcomes

A fresh look at how tools come together shows one spot handling every task without hiccups. Picture after picture reveals what each screen does inside the setup.

1. Navigation Hub

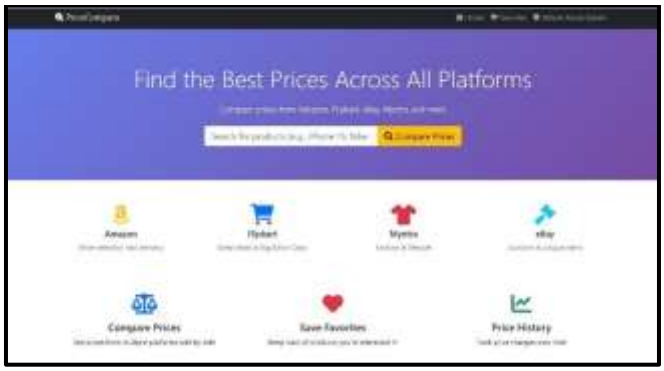

Figure 2 : Centralized Navigation Hub

Starting at the top, everything begins on one main view where choices spread out clearly. Moving through tools feels smooth because they sit in separate blocks with space around them. One path leads to checking prices across stores, another opens controls for spotting refund misuse. There you also find personal spots like saved picks tucked neatly aside. Shifting from shopping help to safety checks happens without jumps or confusion

2. Administrative Command Center

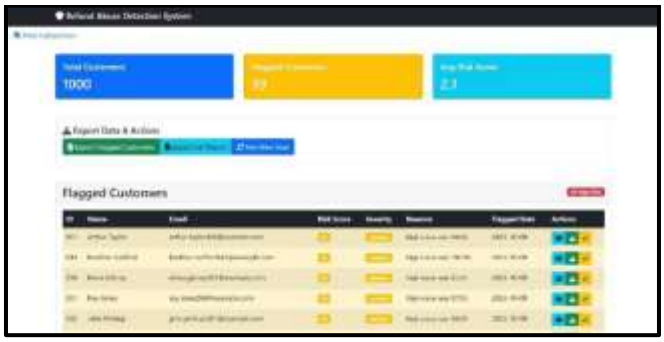

Figure 3 : Refund Abuse Detection System Dashboard

Right off, the admin screen shows a big-picture look at how risky things are across the system. Peeking into numbers, it follows main metrics like total users hitting one thousand, thirty-nine marked cases, along with an average danger mark sitting at two point one. Down below, the list labeled "Flagged Customers" runs on RRP rules, breaking down fine details - say, return rates jumping from half up to nearly eighty-eight percent - sparking auto warnings while letting staff pull out records when needed

3. Consumer Engagement Portal

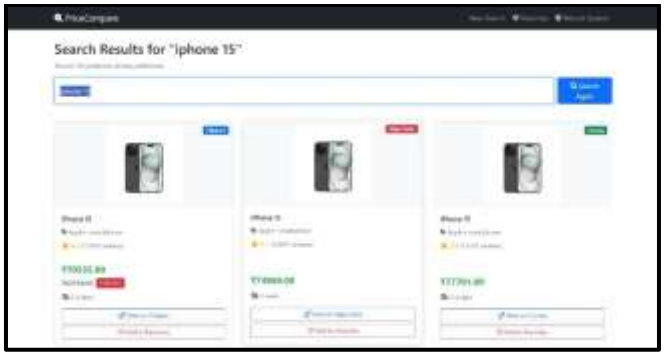

Figure 4 : Price Compare Search Landing Page

Out of the scraping engine comes a stream of live data pulled together in one place. When someone looks up something like iPhone 15, prices pop up instantly - cleaned and sorted - from stores including Flipkart, Vijay Sales, Croma. Ratings sit beside those numbers, discounts highlighted just below. Seeing everything lined up helps spot the strongest offer without flipping through sites. Behind it all, each click gets stored quietly, shaping how future risks are judged according to what that person tends to buy.

V. Conclusion

Not every security move needs alarms. Instead of chasing fraud after it happens, watching how users act reveals more. A mix of smart models - one sharp on patterns, another tuned to subtle shifts - works behind the scenes. This combo spots fake returns using a clear number: how often someone asks for money back. At the same time, shoppers get help checking prices elsewhere. One piece feeds the other while overseeing every move a person makes, not just single clicks or purchases, changes what detection means. Profit stays protected when insight runs deep. Trust grows when friction fades.

Later down the line, computer vision might check how returned products actually look, spotting tricks like worn clothes sent back as new. Instead of just forms, smart summaries built by language models could show investigators clearer pictures of sneaky behavior. Customer service chats may get scanned live, picking up real frustration versus people gaming the rules. These shifts aim to sharpen detection without slowing honest returns. Deeper clues emerge when systems learn to read between the lines. Fraud hides in details most tools miss today.

References

[1] M. Srinivas, M. Kilaru, D. Jain and K. S. Sidhu, "Fraud Detection and Prevention in E-Commerce: Machine Learning Approaches to Secure Transactions," *2025 First International Conference on Advances in Computer Science, Electrical, Electronics, and Communication Technologies (CE2CT)*, Bhimtal, Nainital, India, 2025, pp. 1416-1420, doi: 10.1109/CE2CT64011.2025.10939681.

[2] B. Balasubramanian and K. T. Perannagari, "Investigating Return Behavior in e-commerce: A Data-Driven Study," *2024 IEEE International Conference on Interdisciplinary Approaches in Technology and Management for Social Innovation (IATMSI)*, Gwalior, India, 2024, pp. 1-6, doi: 10.1109/IATMSI60426.2024.10502719.

[3] R. Jhangiani, D. Bein and A. Verma, "Machine Learning Pipeline for Fraud Detection and Prevention in E-Commerce Transactions," *2019 IEEE 10th Annual Ubiquitous Computing, Electronics & Mobile Communication Conference (UEMCON)*, New York, NY, USA, 2019, pp. 0135-0140, doi: 10.1109/UEMCON47517.2019.8992993.

[4] R. Cantini, F. Marozzo, A. Orsino, M. Passarelli and P. Trunfio, "A visual tool for reducing returns in e-commerce platforms," *2021 IEEE 6th International Forum on Research and Technology for Society and Industry (RTSI)*, Naples, Italy, 2021, pp. 474-479, doi: `10.1109/RTSI50628.2021.9597230.

[5] G. T, V. Sathya, S. Vinodhini, E. Chandralekha, V. Usha and R. S, "Improving E-commerce Fraud Detection via Machine Learning: Comparative Evaluation of Model Effectiveness," *2023 International Conference on Innovative Computing, Intelligent Communication and Smart Electrical Systems (ICSES)*, Chennai, India, 2023, pp. 1-6, doi: 10.1109/ICSES60034.2023.10465515.

[6] G. Manoharan, S. D. N. H. Ali, D. M. Sathe, A. Karthik, A. Nagpal and A. Sidana, "Fraud Detection in E-commerce Transactions: A Machine

Learning Perspective," *2024 International Conference on Advances in Computing, Communication and Applied Informatics (ACCAI)*, Chennai, India, 2024, pp. 1-5, doi: 10.1109/ACCAI61061.2024.10601813.

[7] F. Khaled Alarfaj and S. Shahzadi, "Enhancing Fraud Detection in Banking With Deep Learning: Graph Neural Networks and Autoencoders for Real-Time Credit Card Fraud Prevention," in *IEEE Access*, vol. 13, pp. 20633-20646, 2025, doi: 10.1109/ACCESS.2024.3466288

[8] H. S. Alsagri, "Hybrid Machine Learning-Based Multi-Stage Framework for Detection of Credit Card Anomalies and Fraud," in *IEEE Access*, vol. 13, pp. 77039-77048, 2025, doi: 10.1109/ACCESS.2025.3565612.

[9] H. A. Wahshat, R. Al-Rousan, W. A. Al-Haithami, S. I. Ahmad Saany, F. S. Mohamad and R. Kumar, "Current Problems and Proposed Solutions for the Electronic Commerce Sector," *2023 International Conference on Computer Science and Emerging Technologies (CSET)*, Bangalore, India, 2023, pp. 1-6, doi: 10.1109/CSET58993.2023.10346777.

[10] A. Mutemi and F. Bacao, "E-Commerce Fraud Detection Based on Machine Learning Techniques: Systematic Literature Review," in *Big Data Mining and Analytics*, vol. 7, no. 2, pp. 419-444, June 2024, doi: 10.26599/BDMA.2023.9020023.

[11] S. Farsi and M. Chowdhury, "E-Vigilant: An E-Commerce Assistant Smartphone Application with Customer and Seller Fraud Detection, Seller Identity Check, Legal Help, E-Shopping Features," *2024 Second International Conference on Advanced Computing & Communication Technologies (ICACCTech)*, Sonipat, India, 2024, pp. 412-417, doi: 10.1109/ICACCTech65084.2024.00074.

[12] F. Jin, Z. Niu and H. Lang, "A Pattern Based Anti-Fraud Method in C2C Ecommerce Environment," *2010 International Conference on E-Business and E-Government*, Guangzhou, China, 2010, pp. 242-244, doi: 10.1109/ICEE.2010.68.

[13] R. Aggarwal, P. K. Sarangi and A. K. Sahoo, "Credit Card Fraud Detection: Analyzing the Performance of Four Machine Learning Models," *2023 International Conference on Disruptive Technologies (ICDT)*, Greater Noida, India, 2023, pp. 650-654, doi: 10.1109/ICDT57929.2023.10150782.

[14] A. Aditi, A. Dubey, A. Mathur and P. Garg, "Credit Card Fraud Detection Using Advanced Machine Learning Techniques," *2022 Fifth International Conference on Computational Intelligence and Communication Technologies (CCICT)*, Sonepat, India, 2022, pp. 56-60, doi: 10.1109/CCiCT56684.2022.00022